

NYC Yellow Taxi Data Platform

Architecture Overview
End-to-End Governed Data Engineering Solution

Author:
Mallikarjun Reddy

<https://github.com/mallireddy0915/AWS-Data-Engineer>

January 2026

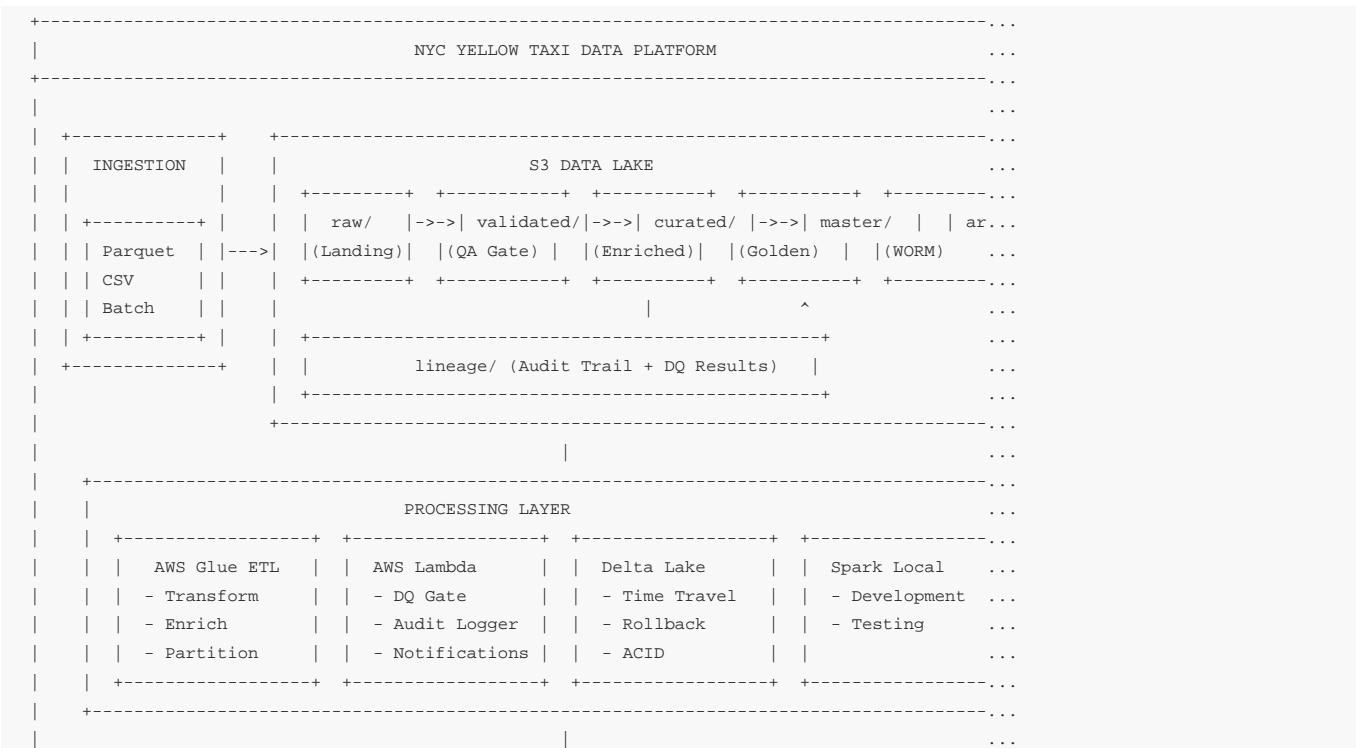
NYC Yellow Taxi Data Platform - Architecture Overview

Executive Summary

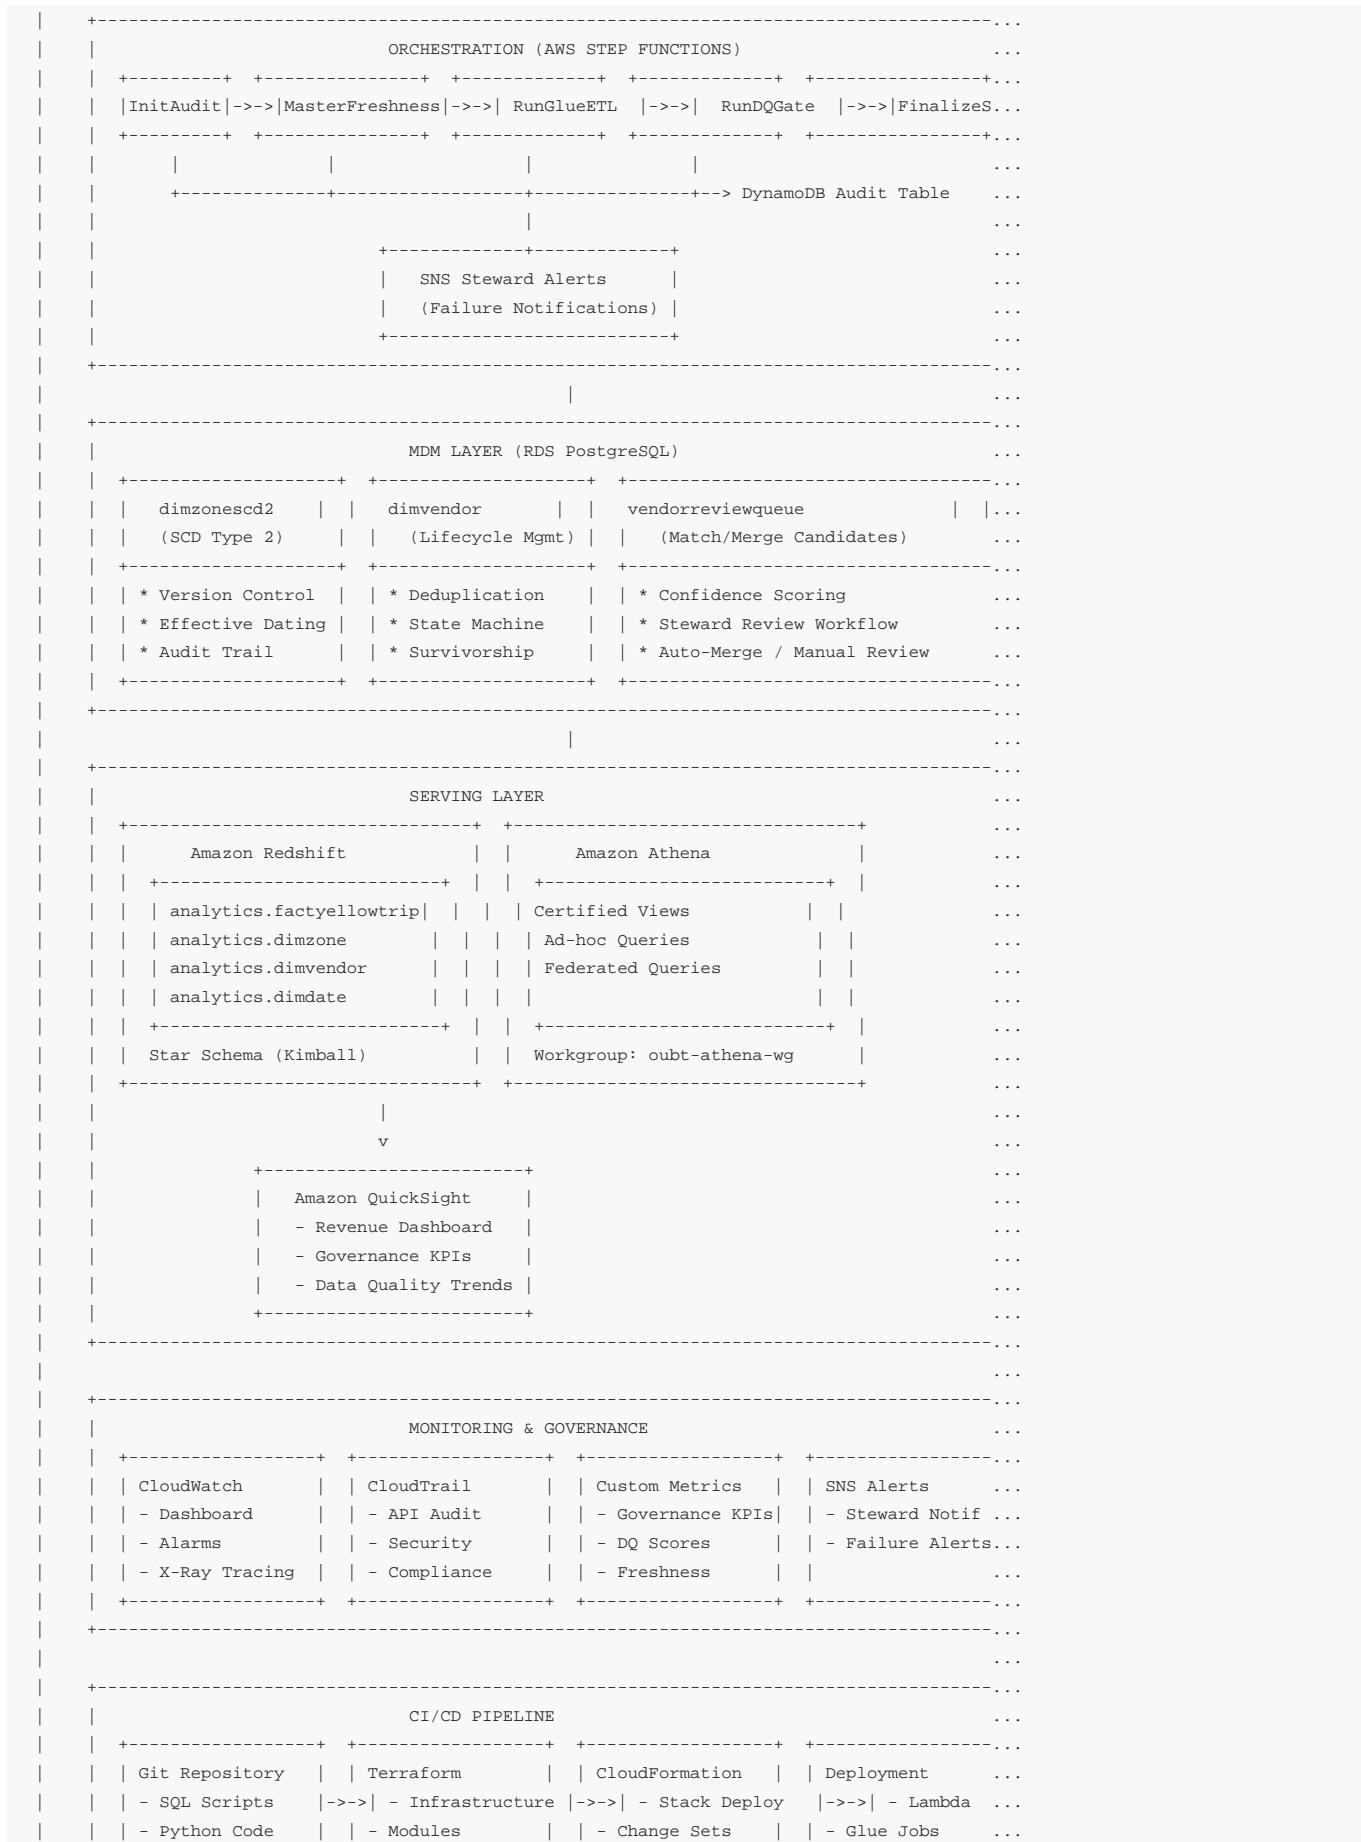
This document presents a comprehensive end-to-end Governed Data Platform built on AWS for NYC Yellow Taxi trip analytics. The platform implements enterprise-grade data engineering patterns including:

- Multi-zone data lake with governance controls
 - Master Data Management (MDM) with deduplication and SCD Type 2
 - Automated orchestration via AWS Step Functions
 - Data quality gates with Great Expectations and AWS Glue DQ
 - Dimensional analytics on Amazon Redshift and Athena
 - Real-time monitoring with CloudWatch dashboards
 - CI/CD deployment via CloudFormation
-

Architecture Diagram



NYC Yellow Taxi Data Platform - Architecture Overview



NYC Yellow Taxi Data Platform - Architecture Overview

```
|   |   | - Config Files   |   | - Environments   |   | - Rollback   |   | - Step Functions...
|   +-----+ +-----+ +-----+ +-----+ ...
|   ...
|   ...
+-----+ ...
```

1. Ingestion Layer (Batch via S3)

Overview

Data arrives as batch files (Parquet, CSV) into the S3 landing zone.

Implementation Details

Component	Description
Source Files	NYC Yellow Taxi trip records (Parquet), Tax..
Landing Zone	s3://arjun-s3-776312084600/raw/
File Formats	Parquet (trips), CSV (reference data)
Ingestion Pattern	Batch upload via AWS CLI, boto3, or schedul..

Key Scripts

- data_ops.py - S3 upload/download operations
- s3_infra.py - Bucket creation and lifecycle configuration

2. Storage Layer - S3 Data Lake with Zones

Zone Architecture

```
+-----+
|          S3 DATA LAKE ZONES          |
+-----+-----+-----+-----+-----+
| raw/      | validated/   | curated/    | master/     | archive/    |
| (Landing) | (QA Gate)  | (Enriched)  | (Golden)   | (WORM)    |
+-----+-----+-----+-----+-----+
| * Immutable | * Quality   | * Business   | * Golden    | * Compliance |
| * Original  | checks       | rules        | records    | retention   |
| * Source of | passed       | applied      | * Strict    | * Immutable  |
| truth       | * Schema     | * MDM        | governance | * 7+ years   |
| * Open       | validated    | enriched    | * Steward   | * Legal hold |
| ingestion   | * Auto-gated | * Analytics  | approval   | enabled     |
|             |               | ready       | required   |           |
+-----+-----+-----+-----+-----+
```

NYC Yellow Taxi Data Platform - Architecture Overview

+-----+-----+
lineage/
(Audit Trail + DQ Results)
* Data provenance tracking
* Transformation history
* Governance audit logs
+-----+-----+

Bucket Lifecycle Configuration

Zone	Retention	Transition	Access Control
raw/	90 days hot	-> Glacier after 90d	Open ingestion
validated/	30 days hot	Archive after 30d	Auto-gated
curated/	1 year	Standard-IA after 60d	Governed transforms
master/	Permanent	None	Strict access control
archive/	7+ years	Deep Archive	Legal hold, WORM

3. Processing Layer - Spark/Glue/Lambda with Delta Lake

AWS Glue ETL

```
# Key transformation logic (day7gluetaxicurated.py)
# 1. Quality gates (validated zone rules)
tripsq = trips.filter(col("tpeppickupdatetime").isNotNull()) \
    .filter(col("tpepdropoffdatetime").isNotNull()) \
    .filter(col("tpeppickupdatetime") <= col("tpepdropoffdatetime")) \
    .filter(col("tripdistance") >= 0) \
    .filter(col("totalamount") >= 0)

# 2. Add partitions
tripsq = tripsq.withColumn("year", year(col("tpeppickupdatetime"))) \
    .withColumn("month", month(col("tpeppickupdatetime")))

# 3. Enrich with zone master data
tripsq = tripsq.join(puzones, "PULocationID", "left") \
    .join(dozones, "DOLocationID", "left")
```

AWS Lambda Functions

Function	Purpose
day11auditlogger	Records pipeline execution to DynamoDB
day11dqgate	Executes Glue Data Quality rulesets
day11masterfreshness	Validates MDM data freshness before ETL
day11notifysteward	Sends SNS alerts on failures

Delta Lake Features

- Time Travel: Query historical data versions
 - ACID Transactions: Consistent reads during writes
 - Rollback/Restore: day15deltarollback_restore.py
-

4. Transformation Layer - SQL Scripts with Version Control

SQL Script Organization

```
sql/
--- 00admin/          # Admin DDL
--- 10transforms/     # Core transformations
--- 20views/          # Analytical views
--- 30quality/        # DQ test queries
--- 40tests/          # Unit tests for SQL
--- 50scd/            # SCD Type 2 logic
--- 51_versioning/    # Version control metadata
--- athena/           # Athena-specific SQL
--- redshift/         # Redshift star schema
```

Script Header Standard (Version Control)

```
-- =====
-- Script:      day4mdmschema.sql
-- Author:      data_engineer
-- Created:    2026-01-15
-- Modified:   2026-01-20
-- Version:    1.2.0
-- Domain:     MDM
-- Description: Master data dimension tables with audit columns
-- =====
```

Key SQL Artifacts

Script	Purpose
day4mdmschema.sql	MDM dimension tables with versioning
510procapplyzonescd2.sql	SCD Type 2 stored procedure
day16starschema.sql	Redshift dimensional model
day17dashboardkpis_view.sql	Governance KPI views

5. Serving Layer - Redshift + Athena

Redshift Star Schema

ANALYTICS STAR SCHEMA			
	dimdate		
	datesk (PK)		
	fulldate		
	year, month, day		
dimzone		factyellowtrip	dimvendor
zonesk (PK)	?-- tripid (PK)	--? vendorsk	
locationid	pickups	vendorid	
borough	dropoffs	vendorname	
zone	datesk (FK)	lifecycle	
servicezone	vendorsk (FK)	state	
mdmversion	puzonesk (FK)		
mdmapprovedby	dozonesk (FK)		
	passengercount		
	tripdistance		
	fareamount		
	totalamount		

Redshift Configuration

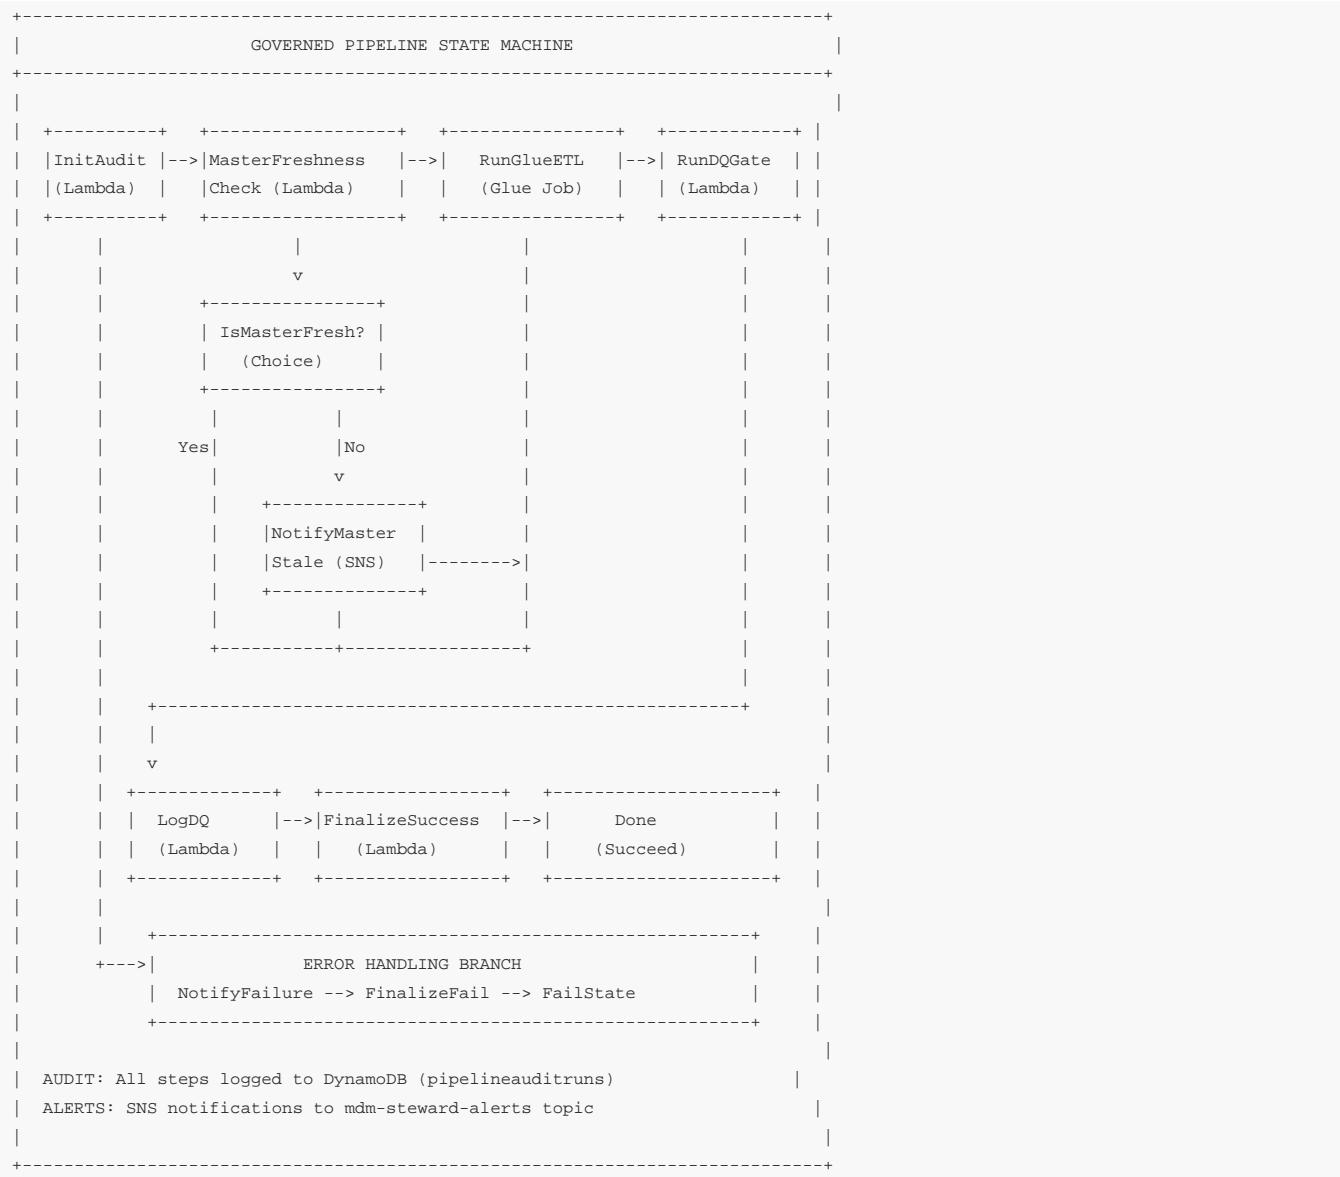
- Cluster: oobt-redshift (ra3.xlplus, single-node demo)
- Distribution: AUTO for fact, ALL for dimensions
- Sort Keys: pickup_ts on fact for time-series queries
- IAM Role: S3 COPY access for data loading

Athena Configuration

- Workgroup: oubt-athena-wg
- Catalog: AWS Glue Data Catalog
- Use Cases: Ad-hoc queries, federated queries, certified views

6. Orchestration - AWS Step Functions

Pipeline State Machine

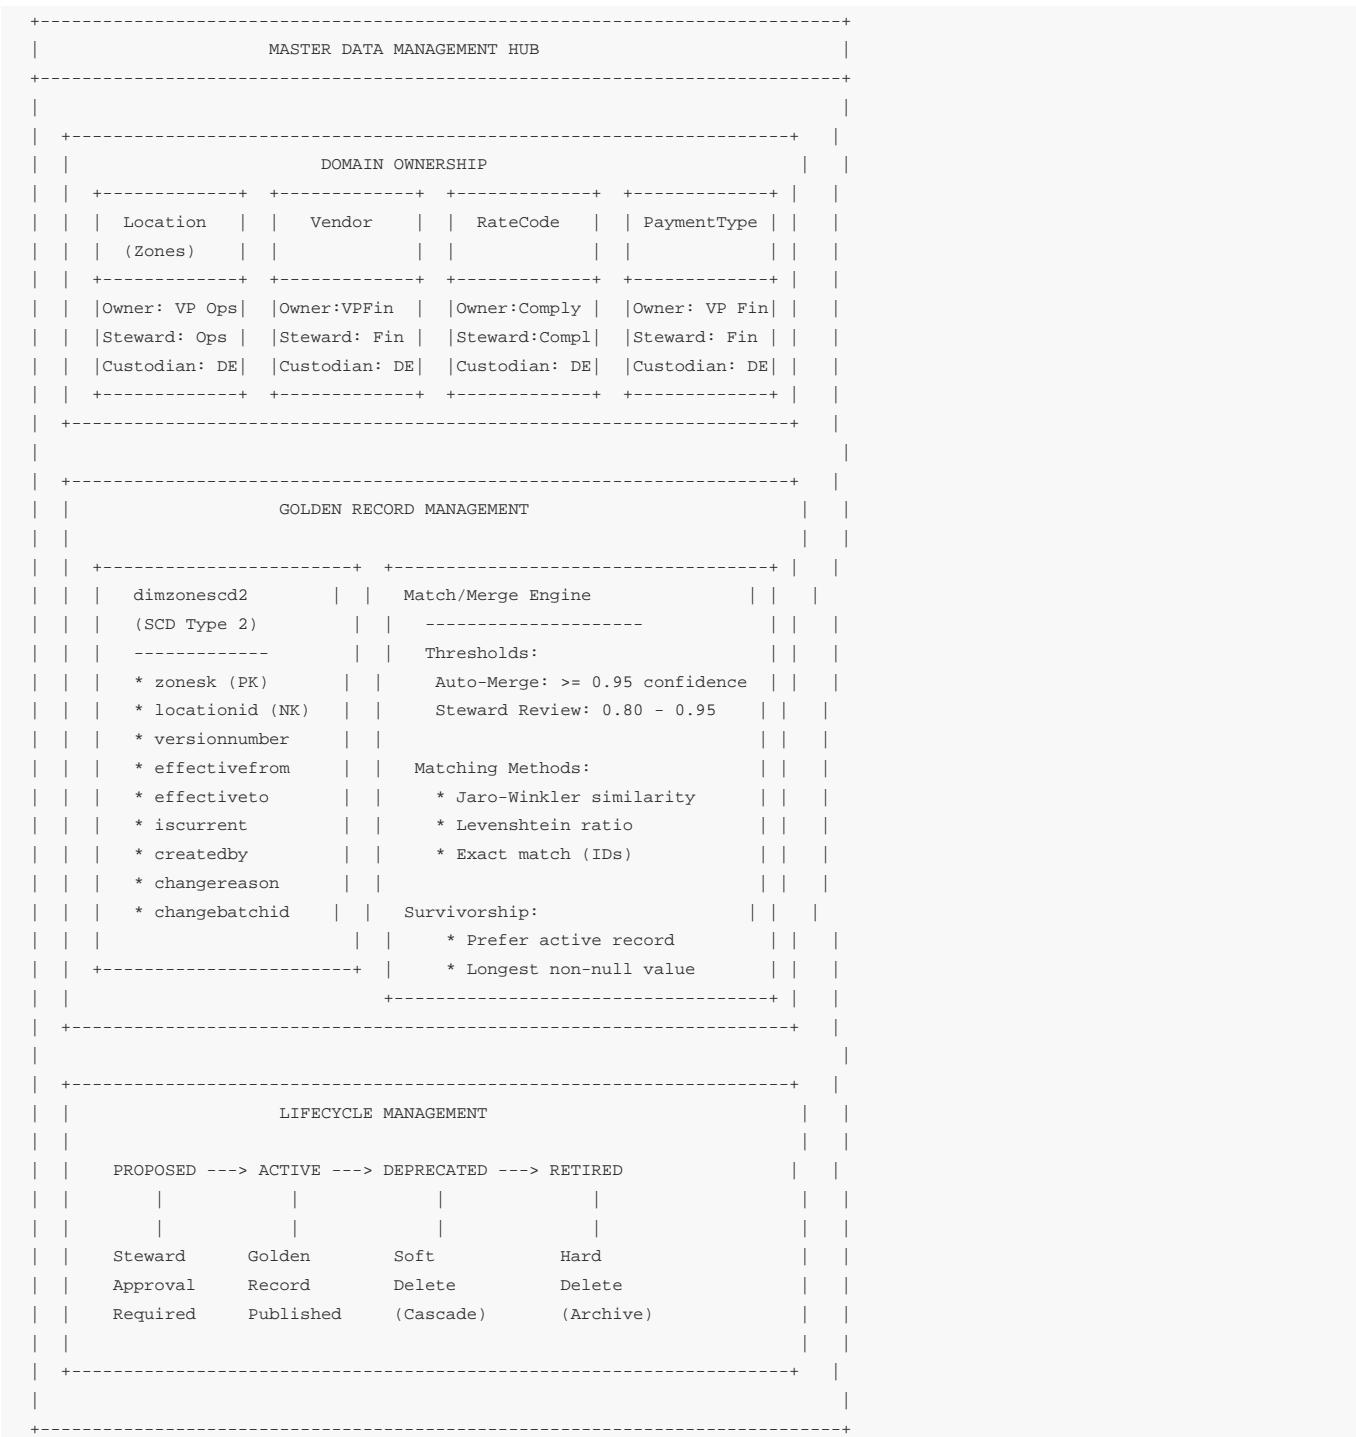


Step Functions Features

- Retry Logic: 3 retries with exponential backoff
- Error Handling: Catch all errors, notify stewards
- Audit Trail: Every step logged to DynamoDB
- State Persistence: Execution history for debugging

7. Master Data Management (MDM) Layer

MDM Architecture



SCD Type 2 Implementation

```
-- Stored procedure: applyzonescd2fromstage
-- Key logic:
```

NYC Yellow Taxi Data Platform - Architecture Overview

```
IF vchanged THEN
    -- Expire current row
    UPDATE mdm.dimzonescd2
    SET effectiveto = vnow, iscurrent = FALSE
    WHERE zonesk = vcurrent.zonesk;
    -- Insert new version
    INSERT INTO mdm.dimzonescd2(
        locationid, borough, zone, servicezone,
        versionnumber, effectivefrom, iscurrent,
        createdby, changereason, changebatchid
    ) VALUES (...);
    -- Audit trail
    INSERT INTO mdm.masterversionaudit(...);
END IF;
```

Match/Merge Rules (vendormatchrules.yaml)

```
thresholds:
  automerge: 0.95      # Automatic merge
  stewardreview: 0.80  # Human review required
fields:
  - name: vendorname
    weight: 0.85
    method: "stringsimilarity"
  - name: vendorid
    weight: 0.15
    method: "exact"
survivorship:
  automergestrategy: "preferactiverecord"
fieldrules:
  vendorname: "longestnon_null"
```

8. Monitoring Layer - CloudWatch Dashboards

Observability Dashboard

OUBT OBSERVABILITY DASHBOARD		
Domain: NYC_Taxi	Owner: Operations Team	Certified: Yes
STEP FUNCTIONS EXECUTIONS	GOVERNANCE KPIs (Custom Metrics)	
-----	-----	
???????????????????????	Pipeline Success Rate: 98.5%	
???????????????????????	Master Freshness: 2.3 hours	
???????????????????????	Orphan Rate PU: 0.02%	
-----	Orphan Rate DO: 0.01%	
# Started # Succeeded	-----	
# Failed # TimedOut		
DATA QUALITY SCORECARD		

NYC Yellow Taxi Data Platform - Architecture Overview

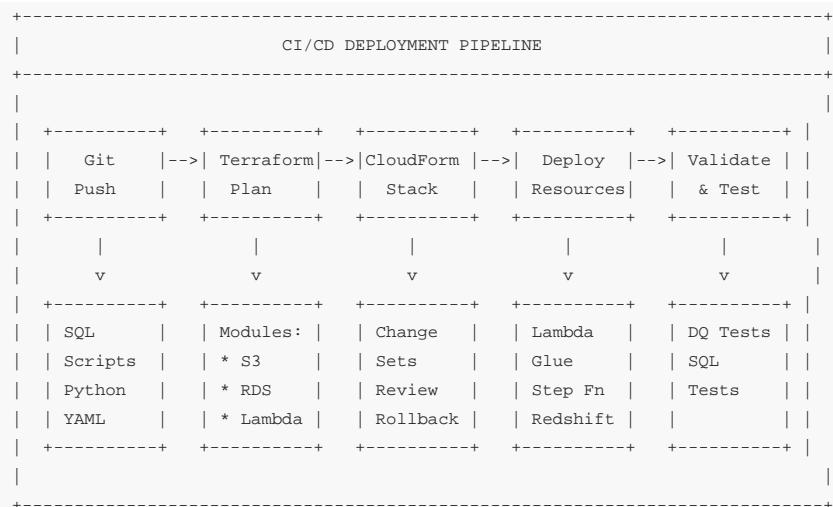
Dimension	Metric	Threshold	Status	
Completeness	Non-null timestamps	>= 99.9%	? PASS	
Validity	Distances >= 0	100%	? PASS	
Consistency	Known zones (no orphan)	<= 0.1% orphan	? PASS	
Timeliness	Curated within 24h	< 24 hours	? PASS	

CloudWatch Alarms

Alarm	Condition	Action
Pipeline Failure	ExecutionsFailed > 0	SNS -> Steward email
DQ Gate Failure	DQ score < threshold	Block pipeline, notify
Master Stale	Freshness > 24h	Alert for review
Lambda Errors	Error count > 5/5min	PagerDuty integration

9. CI/CD Pipeline - CloudFormation

Deployment Architecture



Terraform Modules

```

infra/terraform/
+-- envs/
|   +-- dev.tf
|   +-- prod.tf
|   +-- variables.tfvars
+-- modules/
    +-- lambda/
        +-- main.tf
    
```

NYC Yellow Taxi Data Platform - Architecture Overview

```
|   +-+ variables.tf  
+- rds/  
|   +-+ main.tf  
|   +-+ outputs.tf  
|   +-+ variables.tf  
+- s3/  
  +-+ main.tf  
  +-+ outputs.tf  
  +-+ variables.tf
```

Infrastructure Components (Terraform)

Resource	Configuration
S3 Buckets	Versioning, lifecycle, encryption
RDS PostgreSQL	MDM hub, private subnet
Redshift	ra3.xlplus, S3 COPY role
Lambda Functions	Python 3.11, VPC access
Step Functions	Pipeline orchestration
IAM Roles	Least privilege policies

10. Live Demo Guide (20 Minutes)

Demo 1: Master Data Operations (5 min)

```
# 1. Show current MDM state  
psql -c "SELECT * FROM mdm.dimzone WHERE locationid = 1;"  
# 2. Stage an update  
INSERT INTO mdm.stgzoneupdates (locationid, borough, zone, servicezone, changereason)  
VALUES (1, 'Manhattan', 'Newark Airport Updated', 'EWR', 'Demo update');  
# 3. Apply SCD2 procedure  
CALL mdm.applyzonescd2fromstage('demouser', 'Live demo', 'Manual');  
# 4. Show version history  
SELECT zonesk, locationid, zone, versionnumber, iscurrent, effectivefrom  
FROM mdm.dimzonescd2 WHERE locationid = 1 ORDER BY version_number;
```

Demo 2: Data Quality Monitoring (4 min)

```
# 1. Show Great Expectations suite  
cat governance/greatexpectations/yellowtripssuite.json  
# 2. Run DQ evaluation  
python scripts/day8rundqeveluation.py  
# 3. Show results in CloudWatch  
aws cloudwatch get-metric-data --metric-name "DQRulesPassed" ...  
# 4. Show quality scorecard  
cat docs/quality_scorecard.md
```

Demo 3: Analytics Dashboard (3 min)

NYC Yellow Taxi Data Platform - Architecture Overview

```
-- Connect to Redshift
-- 1. Show star schema
SELECT  FROM analytics.factyellowtrip LIMIT 10;
SELECT  FROM analytics.dimzone LIMIT 10;
-- 2. Run analytics query
SELECT
    dz.borough,
    COUNT(*) as trips,
    SUM(f.totalamount) as revenue
FROM analytics.factyellowtrip f
JOIN analytics.dimzone dz ON f.puzonesk = dz.zonesk
GROUP BY dz.borough
ORDER BY revenue DESC;
```

Demo 4: Batch ETL with Step Functions (4 min)

```
# 1. Start pipeline execution
python scripts/day1startexecutionmanual.py
# 2. Monitor in AWS Console (Step Functions)
# Show state machine graph, execution progress
# 3. Check DynamoDB audit table
aws dynamodb scan --table-name pipelineaudit_runs
# 4. Show SNS notification (if failure)
```

Demo 5: CI/CD Deployment (4 min)

```
# 1. Show Terraform plan
cd infra/terraform
terraform plan
# 2. Show module structure
ls -la modules/
# 3. Deploy change (Lambda update)
terraform apply -auto-approve
# 4. Verify deployment
aws lambda get-function --function-name day1auditlogger
```

Summary Table

Layer	Technology	Key Features
Ingestion	S3, boto3	Batch upload, Parquet/CSV
Storage	S3 Data Lake	5 zones (raw -> master), lif..
Processing	Glue, Lambda, Spark	ETL, DQ gates, Delta Lake
Transformation	SQL (versioned)	SCD2, quality tests
Serving	Redshift, Athena	Star schema, certified views
Orchestration	Step Functions	Governed pipeline, retry/catch
MDM	PostgreSQL	Dedup, match/merge, SCD2
Monitoring	CloudWatch	Dashboards, alarms, X-Ray
CI/CD	Terraform, CloudFormation	IaC, modular deployment

Appendix: Key File References

Category	File Path
Architecture Doc	docs/day20_architecture.md
MDM Architecture	docs/day19mdmarchitecture.md
Lake Zone Design	docs/day6datalakearchitecture.svg
Governance Charter	docs/governance_charter.md
Step Functions	stepfunctions/day11_pipeline.asl.json
Glue ETL	gluejobs/day7gluetaxicurated.py
MDM Schema	sql/day4mdmschema.sql
SCD2 Procedure	sql/50scd/510procapplyzone_scd2.sql
Redshift Schema	sql/redshift/day16starschema.sql
Match Rules	governance/mdm/vendormatchrules.yaml
DQ Suite	governance/greatexpectations/yellowtrips_su..
Dashboard	dashboards/day18/oobtobservabilitydashboard..
Terraform	infra/terraform/modules/

Document Version: 1.0.0

Last Updated: 2026-01-21

Author: Data Engineering Team